

Characterizing transcriptomes using ngs data

T. Källman


BILS/Scilife Lab/Uppsala University

Sep. 2015


Outline

- 1 The transcriptome
- 2 RNA sequence technologies
- 3 RNA-seq analysis
 - Mapping based approach
 - Gene expression from RNA-seq
 - de-novo assembly

The Central Dogma



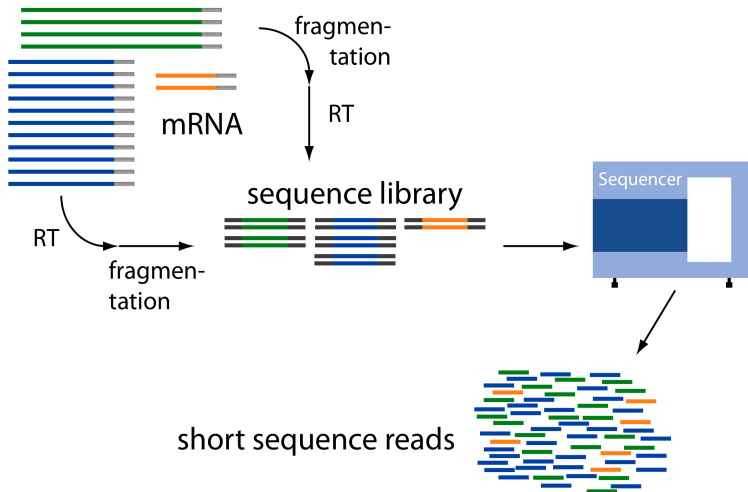
A more complex view



Transcriptomes vs genomes

- Dynamic, not the same over tissues and time points
- Smaller sequence space
- Less repetitive (but large gene families can be found)
- Fairly stable in size? (*eg.* 2-4 fold change among eukaryotes, whereas genome size can vary 1000-fold)
- Genes are often expressed in multiple different splice-variants
- RNA often from only one strand

NGS data



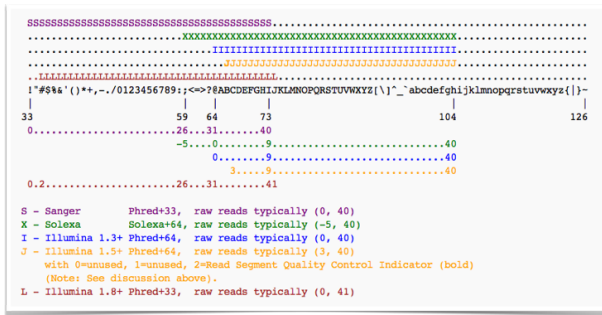


Machine output

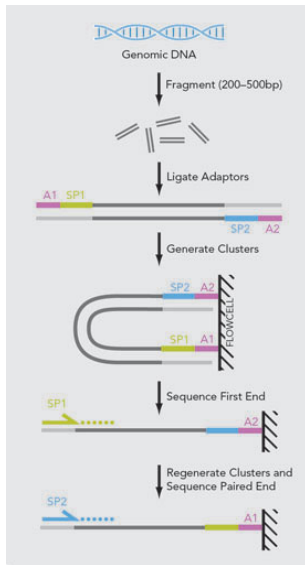
```
@SRR038845.3 HWI-EAS038:6:1:0:1938 length=36
CAACGAGTTCACACCTTGGCCGACAGGCCCGGGTAA
+SRR038845.3 HWI-EAS038:6:1:0:1938 length=36
BA@7>B=>:>>7@7@>>9=BAA?;>52;>:9=8.=A
@SRR038845.41 HWI-EAS038:6:1:0:1474 length=36
CCAATGATTTTTTTCCGTGTTTCAGAATACGGTTAA
+SRR038845.41 HWI-EAS038:6:1:0:1474 length=36
BCCBA@BB@BBBBBAB@B9B@=BABA@A:@693:@B=
@SRR038845.53 HWI-EAS038:6:1:1:360 length=36
GTTCAAAAAGAACTAAATTGTGTCAATAGAAAACCTC
+SRR038845.53 HWI-EAS038:6:1:1:360 length=36
BBCBBBBBBB@@BAB?BBBBCBC>BBBAA8>BBBAA@
```


Sequence quality

- Phred quality scores: $Q = -10 \times \log P$ (High Q = high probability of the base being correct)
- A Phred quality score of 20 to a base, means that the base is called incorrectly in 1 out of 100 times.



Pair-end (PE) sequencing



Pair-end reads


File format

- Two files are created
- The order in files identical and naming of reads are the same with the exception of the end
- The way of naming reads are changing over time so the read names depend on software version


```
@61DFRAAXX100204:1:100:10494:3070/1
AAACAACAGGGCACATTGTCACCTTGTATTTGAAAAACACTTCCGGCCAT
+
ACCCCCCCCCCCCCCCCCCCCCCCCCCCCCBC?CCCCCCCC@@CACCCCCA
```

```
@61DFRAAXX100204:1:100:10494:3070/2
ATCCAAGTTAAACAGAGGCCTGTGACAGACTCTTGGCCCATCGTGTTGATA
+
_^_a^cccegcgghhgZc`ghhc^egggd^_[d]defcdfd^Z^0XWaq^ad
```

Pair-end data




Stranded or not




Basic quality control of raw reads

● FastQC



Basic quality control of raw reads


● FastQC



Basic quality control of raw reads

- RNA-seq is not random sample from the genome eg. GC content might be different
- Highly expressed genes can be frequent and create warnings in quality controls that assumes whole genome data
- Random hexamer in cDNA synthesis might create 'biases' in base frequencies in the beginning of reads

Two main routes for analysis



Haas & Zody (2010), Nature Biotechnology 28, 421–423


Aligning short reads from RNA to genomes

- If available map to the genome sequence
- If no genome sequence one can also map to transcriptome reference
- Make use of available genome annotation (GTF, GFF, BED files)

Galaxy										Analyze Data	Workflow	Shared Data	Visualization	Help	User	Using 0 bytes
Seqname	Source	Feature	Start	End	Score	Strand	Frame	Attributes								
chr12	unknown	exon	87984	88017	.	+	.	gene_id "LOC100288778"; gene_name "LOC100288778"; transcript_id "NR_028269"; tss_id "TSS8200";								
chr12	unknown	exon	88257	88392	.	+	.	gene_id "LOC100288778"; gene_name "LOC100288778"; transcript_id "NR_028269"; tss_id "TSS8200";								
chr12	unknown	exon	88570	88771	.	+	.	gene_id "LOC100288778"; gene_name "LOC100288778"; transcript_id "NR_028269"; tss_id "TSS8200";								
chr12	unknown	exon	88860	89018	.	+	.	gene_id "LOC100288778"; gene_name "LOC100288778"; transcript_id "NR_028269"; tss_id "TSS8200";								
chr12	unknown	exon	89675	89827	.	+	.	gene_id "LOC100288778"; gene_name "LOC100288778"; transcript_id "NR_028269"; tss_id "TSS8200";								
chr12	unknown	exon	90587	90655	.	+	.	gene_id "LOC100288778"; gene_name "LOC100288778"; transcript_id "NR_028269"; tss_id "TSS8200";								
chr12	unknown	exon	90796	91263	.	+	.	gene_id "LOC100288778"; gene_name "LOC100288778"; transcript_id "NR_028269"; tss_id "TSS8200";								
chr12	unknown	exon	147946	148509	.	-	.	gene_id "FAM1380"; gene_name "FAM1380"; transcript_id "NR_026823"; tss_id "TSS11862";								
chr12	unknown	exon	148612	148814	.	-	.	gene_id "FAM1380"; gene_name "FAM1380"; transcript_id "NR_026823"; tss_id "TSS11862";								
chr12	unknown	exon	149052	149412	.	-	.	gene_id "FAM1380"; gene_name "FAM1380"; transcript_id "NR_026823"; tss_id "TSS11862";								
chr12	unknown	CDS	176049	176602	.	+	0	gene_id "IQSEC3"; gene_name "IQSEC3"; p_id "P5442"; transcript_id "NM_001170738"; tss_id "TSS17433";								
chr12	unknown	exon	176049	176602	.	+	.	gene_id "IQSEC3"; gene_name "IQSEC3"; p_id "P5442"; transcript_id "NM_001170738"; tss_id "TSS17433";								
chr12	unknown	start_codon	176049	176051	.	+	.	gene_id "IQSEC3"; gene_name "IQSEC3"; p_id "P5442"; transcript_id "NM_001170738"; tss_id "TSS17433";								
chr12	unknown	exon	186542	186878	.	+	.	gene_id "IQSEC3"; gene_name "IQSEC3"; p_id "P13619"; transcript_id "NM_015232"; tss_id "TSS12565";								
chr12	unknown	CDS	208312	208380	.	+	1	gene_id "IQSEC3"; gene_name "IQSEC3"; p_id "P5442"; transcript_id "NM_001170738"; tss_id "TSS17433";								
chr12	unknown	exon	208312	208380	.	+	.	gene_id "IQSEC3"; gene_name "IQSEC3"; p_id "P13619"; transcript_id "NM_015232"; tss_id "TSS12565";								
chr12	unknown	exon	208312	208380	.	+	.	gene_id "IQSEC3"; gene_name "IQSEC3"; p_id "P5442"; transcript_id "NM_001170738"; tss_id "TSS17433";								
chr12	unknown	CDS	234799	235078	.	+	1	gene_id "IQSEC3"; gene_name "IQSEC3"; p_id "P5442"; transcript_id "NM_001170738"; tss_id "TSS17433";								
chr12	unknown	exon	234799	235078	.	+	.	gene_id "IQSEC3"; gene_name "IQSEC3"; p_id "P5442"; transcript_id "NM_001170738"; tss_id "TSS17433";								
chr12	unknown	exon	246577	246793	.	-	.	gene_id "LOC574538"; gene_name "LOC574538"; transcript_id "NR_033859"; tss_id "TSS17153";								
chr12	unknown	CDS	247433	248520	.	+	0	gene_id "IQSEC3"; gene_name "IQSEC3"; p_id "P5442"; transcript_id "NM_001170738"; tss_id "TSS17433";								
chr12	unknown	exon	247433	248520	.	+	.	gene_id "IQSEC3"; gene_name "IQSEC3"; p_id "P13619"; transcript_id "NM_015232"; tss_id "TSS12565";								
chr12	unknown	exon	247433	248520	.	+	.	gene_id "IQSEC3"; gene_name "IQSEC3"; p_id "P5442"; transcript_id "NM_001170738"; tss_id "TSS17433";								
chr12	unknown	CDS	247439	248520	.	+	0	gene_id "IQSEC3"; gene_name "IQSEC3"; p_id "P13619"; transcript_id "NM_015232"; tss_id "TSS12565";								
chr12	unknown	start_codon	247439	247441	.	+	.	gene_id "IQSEC3"; gene_name "IQSEC3"; p_id "P13619"; transcript_id "NM_015232"; tss_id "TSS12565";								

Aligning short reads from RNA to genomes

- Large number of programs available: Star, Tophat, Subread etc
- Important feature: Allow for spliced mapping



Aligning short reads from RNA to genomes

- After mapping perform QC of the output


```
read_distribution.py -i PairedStrandSpecific_51mer_Human_hg19.bam -r hg19.refseq.bed12
```

Output:

Group	Total_bases	Tag_count	Tags/Kb
CDS_Exons	33302033	20002271	600.63
5'UTR_Exons	21717577	4408991	203.01
3'UTR_Exons	15347845	3643326	237.38
Introns	1132597354	6325392	5.58
TSS_up_1kb	17957047	215331	11.99
TSS_up_5kb	81621382	392296	4.81
TSS_up_10kb	149730983	769231	5.14
TES_down_1kb	18298543	266161	14.55
TES_down_5kb	78900674	729997	9.25
TES_down_10kb	140361190	896882	6.39

Example workflow

- Tophat: Aligns reads to genome (allows for spliced read mapping)
- Cufflinks: Extract transcripts from spliced read alignments
- Cuffmerge: Merge results from multiple Cufflinks results
- Cuffdiff: Detect differential gene expression



Trapnell *et al.* (2012), Nature Protocols 7, 562–578

Tophat

- 1 Efficient and fast alignment to the genome using bowtie2
- 2 Create a data base of putative splice junctions from the reads mapping in step 1
- 3 Map reads that did not map in step 1 using the splice information

QC of mapped reads

Reads should mostly map to known (annotated genes)

```
read_distribution.py -i PairedStrandSpecific_5lmer_Human_hg19.bam -r hg19.refseq.bed12
```


Output:

Group	Total_bases	Tag_count	Tags/Kb
CDS_Exons	33302033	20002271	600.63
5'UTR_Exons	21717577	4408991	203.01
3'UTR_Exons	15347845	3643326	237.38
Introns	1132597354	6325392	5.58
TSS_up_1kb	17957047	215331	11.99
TSS_up_5kb	81621382	392296	4.81
TSS_up_10kb	149730983	769231	5.14
TES_down_1kb	18298543	266161	14.55
TES_down_5kb	78900674	729997	9.25
TES_down_10kb	140361190	896882	6.39


QC of mapped reads

Most splice event should be known and canonical (GU-AG)

splicing junctions




splicing events




Cufflinks


a Splice-align reads to the genome



b Build a graph representing alternative splicing events



c Traverse the graph to assemble variants



d Assembled isoforms

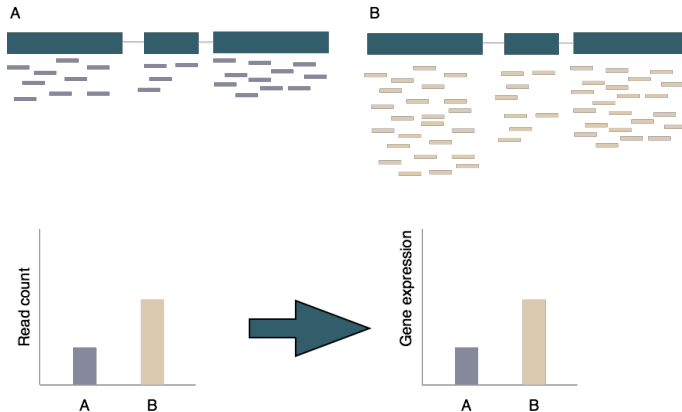


Nature Reviews | Genetics

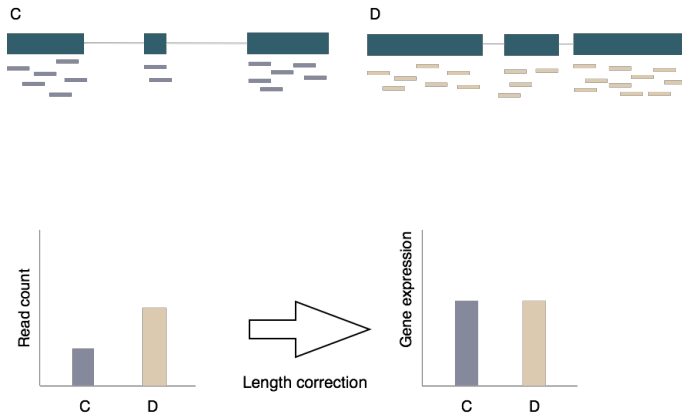
Cuffdiff

- Program that estimate expression levels and identify differentially expressed genes from ngs alignments
- Basically uses the read data to estimate dispersion parameters (the amount of deviation from a Poisson distr.)
- Genes that show patterns deviating from the above expectations are differentially expressed between treatments
- Will work also for detection of isoform differential expression

From counts to gene expression



From counts to gene expression



Not all reads are the same



	union	intersection_strict	intersection_nonempty
	gene_A	gene_A	gene_A
	gene_A	no_feature	gene_A
	gene_A	no_feature	gene_A
	gene_A	gene_A	gene_A
	gene_A	gene_A	gene_A
	ambiguous	gene_A	gene_A
	ambiguous	ambiguous	ambiguous

from: <http://www-huber.embl.de/users/anders/HTSeq/doc/count.html>

Normalized expression Values

- Mapped read counts are normalized for both length of the transcript they map to and total depth of sequencing.
- Count data is hence converted to: Reads/Fragments per kb of transcript length and million mapped reads (RPKM or FPKM)

Experimental design




Experimental design

- Count reads (convert to RPKM/FPKM?)
- Small number of reads (= low RPKM/FPKM values) often non-significant
- Remember that Fold change is not the same as significance

	Condition 1	Condition 2	Fold_Change	Significant?
Gene A	1	2	2-fold	No
Gene B	100	200	2-fold	Yes

Two main routes for analysis



Haas & Zody (2010), Nature Biotechnology 28, 421–423

Major challenges in relation to genome assembly

- Genes show different levels of gene expression, hence uneven coverage among genes
- Many genes are expressed in different isoforms
- As sequence depth increase detected number of loci increase. (What is actually expressed?)
- Sequence error from highly expressed genes might be seen more often than "true" sequences from lowly expressed genes


Several programs available

- SOAP-denovo TRANS
- Oases
- Trans-ABYSS
- Trinity


All of them uses de Bruijn graphs to cope with the data and many of them have been developed from a genome assembly program

Trinity


a Generate all substrings of length k from the reads



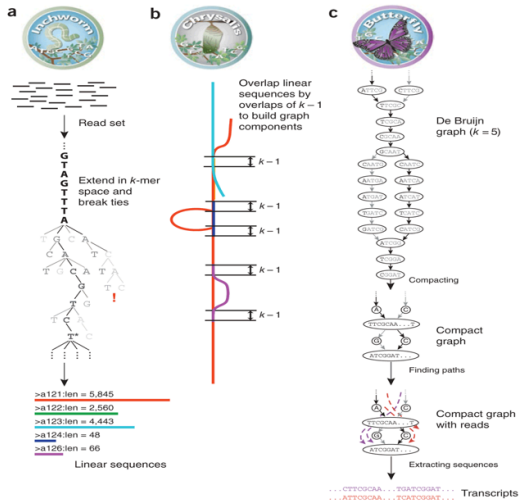
b Generate the De Bruijn graph



c Collapse the De Bruijn graph



Trinity



Summary - with ref.

- Map to genome allow for spliced alignment
- If novel transcripts of interest: use method that can re-create transcripts from mapped reads (Cufflinks, Scripture or Bayesemblem)
NB! In well annotated genomes most reads should map to known genes
- If interest is expression of known genes/exons: Use available annotation for analysis
- Spend time on experimental design and more replicates gives more power in gene expression analysis

Summary - without ref.

- Assemble using your favourite assembler
- Spend lots of time in assessing the results (compare to related species, look for ORFs etc)
- Often large number of partial transcripts (hence often large number of contigs).
- Merge with other data from transcripts?